

Concolic Execution for Fuzzing in Julia

Final project for 6.858 (Spring 2018)

Valentin Churavy

JuliaLab@CSAIL

16. May 2018



- 1 Concolic Execution
- 2 Julia in brief
- 3 Implementation
- 4 Examples
- 5 Limitations

Welcome to Lab 3

Welcome to Lab 3

... but in Julia!

Welcome to Lab 3

... but in Julia!

General strategy

- Generate symbolic trace of the program for a given input

Based on [4, 3]

Welcome to Lab 3

... but in Julia!

General strategy

- Generate symbolic trace of the program for a given input
- Feed symbolic trace to Z3 to generate inputs that cause exploration of unvisited branches

Based on [4, 3]

Welcome to Lab 3

... but in Julia!

General strategy

- Generate symbolic trace of the program for a given input
- Feed symbolic trace to Z3 to generate inputs that cause exploration of unvisited branches
- Iterate until all branches are visited, record which branches throw exceptions

Based on [4, 3]

Julia [1] is a dynamic high-performance language, it looks like Python, runs like Fortran, and talks like Lisp. It has interesting language design properties, like a rich type system, multiple dispatch, macros, and staged programming. It uses a compiler based on LLVM[5].

Julia [1] is a dynamic high-performance language, it looks like Python, runs like Fortran, and talks like Lisp. It has interesting language design properties, like a rich type system, multiple dispatch, macros, and staged programming. It uses a compiler based on LLVM[5].

This is not a Julia tutorial!

I am going to talk about a couple of advanced concepts and use a development version of Julia with experimental tooling!
If you want to learn more about Julia, ask me later!

Julia [1] is a dynamic high-performance language, it looks like Python, runs like Fortran, and talks like Lisp. It has interesting language design properties, like a rich type system, multiple dispatch, macros, and staged programming. It uses a compiler based on LLVM[5].

This is not a Julia tutorial!

I am going to talk about a couple of advanced concepts and use a development version of Julia with experimental tooling!
If you want to learn more about Julia, ask me later!

AD in 5s

Example time!

- Goal: Be able to deal with arbitrary user code

- Goal: Be able to deal with arbitrary user code
- Trace an invocation of a function, e.g. record recursively, which calls occur and which input + output arguments they have.

- Goal: Be able to deal with arbitrary user code
- Trace an invocation of a function, e.g. record recursively, which calls occur and which input + output arguments they have.
- Mark/Taint input arguments as symbolic

- Goal: Be able to deal with arbitrary user code
- Trace an invocation of a function, e.g. record recursively, which calls occur and which input + output arguments they have.
- Mark/Taint input arguments as symbolic
- Issue: The trace is linearized and no longer contains control-flow

- Goal: Be able to deal with arbitrary user code
- Trace an invocation of a function, e.g. record recursively, which calls occur and which input + output arguments they have.
- Mark/Taint input arguments as symbolic
- Issue: The trace is linearized and no longer contains control-flow
- Solution: Rewrite and rejit user-provided functions and insert calls to the assert function, which records conditionals that are used as part of the control-flow.

- Goal: Be able to deal with arbitrary user code
- Trace an invocation of a function, e.g. record recursively, which calls occur and which input + output arguments they have.
- Mark/Taint input arguments as symbolic
- Issue: The trace is linearized and no longer contains control-flow
- Solution: Rewrite and rejit user-provided functions and insert calls to the assert function, which records conditionals that are used as part of the control-flow.
- Filter concolic trace to get a trace that only depends on the tainted arguments

- Goal: Be able to deal with arbitrary user code
- Trace an invocation of a function, e.g. record recursively, which calls occur and which input + output arguments they have.
- Mark/Taint input arguments as symbolic
- Issue: The trace is linearized and no longer contains control-flow
- Solution: Rewrite and rejit user-provided functions and insert calls to the assert function, which records conditionals that are used as part of the control-flow.
- Filter concolic trace to get a trace that only depends on the tainted arguments
- Cut trace at branch and negate branch condition

- Goal: Be able to deal with arbitrary user code
- Trace an invocation of a function, e.g. record recursively, which calls occur and which input + output arguments they have.
- Mark/Taint input arguments as symbolic
- Issue: The trace is linearized and no longer contains control-flow
- Solution: Rewrite and rejit user-provided functions and insert calls to the assert function, which records conditionals that are used as part of the control-flow.
- Filter concolic trace to get a trace that only depends on the tainted arguments
- Cut trace at branch and negate branch condition
- Translate symbolic trace to SMT2/Z3, if satisfiable use the input value to continue exploration

Solving the type-constraint problem with Cassette

One of the challenges was to be able to taint arbitrary user-code, including user code that has type-constraints.

Solving the type-constraint problem with Cassette

One of the challenges was to be able to taint arbitrary user-code, including user code that has type-constraints.

Cassette

- Cassette.jl [6] developed here at MIT provides a framework to extend Julia with non-standard execution models, like AD or symbolic execution, through contextual dispatch

Solving the type-constraint problem with Cassette

One of the challenges was to be able to taint arbitrary user-code, including user code that has type-constraints.

Cassette

- Cassette.jl [6] developed here at MIT provides a framework to extend Julia with non-standard execution models, like AD or symbolic execution, through contextual dispatch
- User-code is executed within a *context* in which the context author can provide alternative *primitives* and can dynamically inject passes into the compiler

Solving the type-constraint problem with Cassette

One of the challenges was to be able to taint arbitrary user-code, including user code that has type-constraints.

Cassette

- Cassette.jl [6] developed here at MIT provides a framework to extend Julia with non-standard execution models, like AD or symbolic execution, through contextual dispatch
- User-code is executed within a *context* in which the context author can provide alternative *primitives* and can dynamically inject passes into the compiler
- Usage for this project

Solving the type-constraint problem with Cassette

One of the challenges was to be able to taint arbitrary user-code, including user code that has type-constraints.

Cassette

- Cassette.jl [6] developed here at MIT provides a framework to extend Julia with non-standard execution models, like AD or symbolic execution, through contextual dispatch
- User-code is executed within a *context* in which the context author can provide alternative *primitives* and can dynamically inject passes into the compiler
- Usage for this project
 - Trace generation through pre-/post-hooks on every function call

Solving the type-constraint problem with Cassette

One of the challenges was to be able to taint arbitrary user-code, including user code that has type-constraints.

Cassette

- Cassette.jl [6] developed here at MIT provides a framework to extend Julia with non-standard execution models, like AD or symbolic execution, through contextual dispatch
- User-code is executed within a *context* in which the context author can provide alternative *primitives* and can dynamically inject passes into the compiler
- Usage for this project
 - Trace generation through pre-/post-hooks on every function call
 - Rewrite IR to inject assert statements

Solving the type-constraint problem with Cassette

One of the challenges was to be able to taint arbitrary user-code, including user code that has type-constraints.

Cassette

- Cassette.jl [6] developed here at MIT provides a framework to extend Julia with non-standard execution models, like AD or symbolic execution, through contextual dispatch
- User-code is executed within a *context* in which the context author can provide alternative *primitives* and can dynamically inject passes into the compiler
- Usage for this project
 - Trace generation through pre-/post-hooks on every function call
 - Rewrite IR to inject assert statements
 - Taint propagation through the metadata system of Cassette

Solving the type-constraint problem with Cassette

One of the challenges was to be able to taint arbitrary user-code, including user code that has type-constraints.

Cassette

- Cassette.jl [6] developed here at MIT provides a framework to extend Julia with non-standard execution models, like AD or symbolic execution, through contextual dispatch
- User-code is executed within a *context* in which the context author can provide alternative *primitives* and can dynamically inject passes into the compiler
- Usage for this project
 - Trace generation through pre-/post-hooks on every function call
 - Rewrite IR to inject assert statements
 - Taint propagation through the metadata system of Cassette
 - Context specific primitives of operations (like addition) to propagate taint or to create taint (rand)

Extensible design: `rand`

One problem in concolic execution is how to deal with random variables, concolic execution depends on being able to *deterministically* execute programs. `rand` is implemented as a primitive that generates taint and the return value of a seen execution of `rand` can be controlled and during fuzzing Z3 will generate values for `rand` much in the same way it does for other inputs.

This could be extended to other foreign-calls like reading from the filesystem.

- Proof some properties
- Fuzzing!
- Loops, structs, random variables

- Reasoning about the type-domain

- Reasoning about the type-domain
- Taint analysis through Arrays

- Reasoning about the type-domain
- Taint analysis through Arrays
- Taint analysis with non-integer inputs

- Reasoning about the type-domain
- Taint analysis through Arrays
- Taint analysis with non-integer inputs
 - Strings

- Reasoning about the type-domain
- Taint analysis through Arrays
- Taint analysis with non-integer inputs
 - Strings
 - Structs

- Reasoning about the type-domain
- Taint analysis through Arrays
- Taint analysis with non-integer inputs
 - Strings
 - Structs
 - Floating-point values

- Reasoning about the type-domain
- Taint analysis through Arrays
- Taint analysis with non-integer inputs
 - Strings
 - Structs
 - Floating-point values
- Arithmetic on Bool

- Using the Klee query language

- Using the Klee query language
- Maybe deeper integration with Klee[2]

- Using the Klee query language
- Maybe deeper integration with Klee[2]
- Only intercept intrinsics and built-ins instead of higher level methods

- Using the Klee query language
- Maybe deeper integration with Klee[2]
- Only intercept intrinsics and built-ins instead of higher level methods
- Better analysis of function to taint usage of global variables and foreigncalls

Where to find this work

<https://github.com/vchuravy/ConcolicFuzzer.jl>

Where to find this work

<https://github.com/vchuravy/ConcolicFuzzer.jl>

If you want to learn more about Julia!

- 1 Write me an e-mail or go to <https://julialang.org>
- 2 Find me in 32-G785



Jeff Bezanson et al. 'Julia: A Fast Dynamic Language for Technical Computing'. In: (Sept. 2012). arXiv: 1209.5145 [cs.PL].



Cristian Cadar, Daniel Dunbar and Dawson Engler. 'KLEE: Unassisted and Automatic Generation of High-Coverage Tests for Complex Systems Programs'. In: (). URL: <http://hci.stanford.edu/cstr/reports/2008-03.pdf>.



Cristian Cadar et al. 'EXE: Automatically Generating Inputs of Death'. In: *Proceedings of the 13th ACM Conference on Computer and Communications Security*. CCS '06. New York, NY, USA: ACM, 2006, pp. 322–335.



Patrice Godefroid, Nils Klarlund and Koushik Sen. 'DART: Directed Automated Random Testing'. In: *SIGPLAN Not.* 40.6 (June 2005), pp. 213–223.



Chris Lattner and Vikram Adve. 'LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation'. In: *Proceedings of the International Symposium on Code Generation and Optimization: Feedback-directed and Runtime Optimization*. CGO '04. Washington, DC, USA: IEEE Computer Society, 2004, pp. 75–.



Jarrett Revels. *Cassette.jl*. 2018. URL: <https://github.com/jrevels/Cassette.jl>.